

Tutorial:

Machine learning ground states with RNN wavefunctions

Navigating Quantum and AI Career Trajectories: A mini-course on
Computational Methods and their Applications

Mohamed Hibat-Allah
University of Waterloo | Perimeter Institute

May 24th, 2024

In this tutorial, you will learn how to use exact diagonalization (ED) and RNN wavefunctions to find the ground state of quantum many-body systems.

Recall from previous lectures in this school that the goal of Variational Monte Carlo (VMC) is to find the ground state of quantum many-body systems, where the main idea is to take an ansatz (called a variational wavefunction) “ $|\Psi_\lambda\rangle$ ” to find the ground state energy E_G by minimizing the variational energy $E_\lambda = \langle \Psi_\lambda | \hat{H} | \Psi_\lambda \rangle$ over a set of the variational parameters $\lambda = \{\lambda_1, \lambda_2, \dots, \lambda_{N_p}\}$. The latter can be done using a learning algorithm such as gradient descent. Here, we will use the learning algorithm known as the Adam optimizer [1].

In this tutorial, we use **positive RNN wavefunctions** to find the ground state of the Transverse-field Ferromagnetic Ising Model (TFIM) in 1D, whose Hamiltonian is given by

$$\hat{H}_{\text{TFIM}} = -J_z \sum_{i=1}^{N-1} \hat{\sigma}_i^z \hat{\sigma}_{i+1}^z - B_x \sum_{i=1}^N \hat{\sigma}_i^x,$$

where N is the number of spins, J_z , B_x are respectively the uniform ferromagnetic coupling between the spins and the uniform transverse magnetic field applied to all the spins. Here we assume open boundary conditions. For the sake of simplicity, we restrict the system size to $N = 10$ spins but you are encouraged to investigate larger system sizes in case you are interested. The parameter values are also restricted to $J_z = 1$ and $B_x = 1$, which corresponds to the critical point of this model. The jupyter notebook of this tutorial can be found on [Google colab](#).

Assuming a given spin configuration $\sigma = (\sigma_1, \sigma_2, \dots, \sigma_N)$ along the z-basis, where $\sigma_i = \pm 1$, we recall from the RNN lectures that the positive RNN wavefunction architecture (shown in Fig. 1) computes the amplitude $\Psi_{\text{RNN}}(\sigma)$ as follows:

- **Step 1:** the green box “RNN” in Fig. 1 takes two inputs $(\sigma_{n-1}, \mathbf{h}_{n-1})$ and outputs \mathbf{h}_n as

$$\mathbf{h}_n = f(U\sigma_{n-1} + V\mathbf{h}_{n-1} + \mathbf{b}),$$

where f is a non-linear activation function (equal to Tanh by default in the [jupyter notebook](#)). The dimension of the memory state \mathbf{h}_n is known as the “**number of memory units**”, and you will be able to examine its effect in this tutorial.

- **Step 2:** \mathbf{h}_n is fed into a Softmax fully-connected layer “S” (shown using magenta circles in Fig. 1) to compute the conditional probability P_n as

$$P_n \equiv P(\sigma_n | \sigma_{n-1}, \dots, \sigma_1) = \text{Softmax}(W\mathbf{h}_n + \mathbf{c}) \cdot \sigma_n,$$

where σ_n is a one-hot representation of the spin σ_n .

- **Step 3:** the amplitude of the spin configuration σ is computed as

$$\Psi_{\text{RNN}}(\sigma) = \sqrt{P_1 P_2 \dots P_N}.$$




Figure 1: An illustration of positive RNN wavefunctions [2].

Within this set-up, the variational parameters λ that we will be optimizing over are the weights U, V, W and the biases \mathbf{b}, \mathbf{c} . You can read more about this approach in Ref. [2].

To verify our results, we compare with exact diagonalization (ED), which can be thought of as a diagonalization of the 1D TFIM Hamiltonian. As an illustration, if $N = 2$, then $\hat{H}_{\text{TFIM}} |\sigma_1 \sigma_2\rangle = -J_z \sigma_1 \cdot \sigma_2 |\sigma_1 \sigma_2\rangle - B_x |\bar{\sigma}_1 \sigma_2\rangle - B_x |\sigma_1 \bar{\sigma}_2\rangle$, where $\bar{\sigma}_i$ denotes that the spin σ_i is flipped. As a consequence, the Hamiltonian can be put in matrix format with size $2^N \times 2^N = 2^2 \times 2^2$ as

$$\hat{H}_{\text{TFIM}} = \begin{bmatrix} -J_z & -B_x & -B_x & 0 \\ -B_x & J_z & 0 & -B_x \\ -B_x & 0 & J_z & -B_x \\ 0 & -B_x & -B_x & -J_z \end{bmatrix},$$

which one can diagonalize to find the exact ground state.

1. Check using exact diagonalization (ED), for $N = 10$, $J_z = 1 = B_x$, that the ground state of 1D TFIM has positive amplitudes.

Hint: You may find it helpful to plot the amplitudes of the ground state provided by the function “ED_1DTFIM” in Sec. 1 of the [jupyter notebook](#). You are more than welcome to check how the diagonalization is performed in the code, but you are not required to understand all the details.

2. The property found in the previous question is a consequence of the negativity of the non-diagonal elements of the TFIM Hamiltonian (cf. Perron-Frobenius theorem). Argue, in this case, that it is enough to use a **positive RNN wavefunction**.
3. Estimate the ground state energy of the 1D TFIM using a positive RNN wavefunction for $N = 10$, $J_z = 1 = B_x$ (see Sec. 2 of the [jupyter notebook](#)). Check the validity of the zero-variance principle (i.e., does the energy variance approach zero?). Compare with the ground state energy given by ED (see Sec. 3 of the [jupyter notebook](#)).
4. Study the effect of a certain hyperparameter:
 - **Hyperparameter 1:** number of memory units (i.e., the size of the hidden state \mathbf{h}_n).
 - **Hyperparameter 2:** activation function f .
 - **Hyperparameter 3:** number of samples.
 - **Hyperparameter 4:** learning rate.
 - **Hyperparameter 5:** type of RNN cell, i.e., either a vanilla or a GRU cell.

Play with your chosen hyperparameter and show how the accuracy changes. You can use a for loop to run multiple VMC optimizations and gather the energy variances in one plot to compare the convergence behavior. Feel free to add more code and make plots that visualize the trends.

For the study of activation functions, you can use Relu, Elu, Tanh, Sigmoid, or no activation function. For the study of the number of memory units and the number of samples, you can try changing the corresponding values in powers of 2. You may want to check Ref. [2] for inspiration.

5. Try switching the transverse magnetic field B_x from +1 to -1. Do you expect the ground state energy to change? (You can check with ED if you are not sure about your answer). Is the ground state still positive? Can the pRNN wavefunction still find an approximate estimation of the ground state energy? Why? What can we do to fix the issue?

PS: If you have finished the tutorial early, you can take the best hyperparameters you found and play with larger system sizes or different values of the transverse magnetic field B_x , to see how the positive RNN wavefunction performs (see Sec. 4 of the [jupyter notebook](#)).

To compare with the exact results, you can make use of the [exact expression of the ground state energy](#) at the critical point ($J_z = 1, B_x = 1$) for open boundary conditions:

$$E_G = 1 - \frac{1}{\sin\left(\frac{\pi}{2(2N+1)}\right)}.$$

If you are interested to see a tutorial about the application of complex RNN wavefunctions, you can check [this jupyter notebook](#) with more details in Ref. [2]. For more details about the applications of RNN wavefunctions you can check Refs. [2–10].

Finally, if you are interested in playing with Transformer wave functions, you can check the code posted on this repository https://github.com/APRIQuOt/VMC_with_LPTFs.

Acknowledgments

I would like to acknowledge Lauren Hayward for her valuable help with the writing and editing of this tutorial. I also thank Sylvain Capponi for the helpful feedback.

References

1. Kingma, D. P. & Ba, J. *Adam: A Method for Stochastic Optimization* 2014. <https://arxiv.org/abs/1412.6980>.
2. Hibat-Allah, M., Ganahl, M., Hayward, L. E., Melko, R. G. & Carrasquilla, J. Recurrent neural network wave functions. *Phys. Rev. Research* **2**, 023358. <https://link.aps.org/doi/10.1103/PhysRevResearch.2.023358> (2 June 2020).
3. Roth, C. *Iterative Retraining of Quantum Spin Models Using Recurrent Neural Networks* 2020. <https://arxiv.org/abs/2003.06228>.
4. Carrasquilla, J. & Torlai, G. How To Use Neural Networks To Investigate Quantum Many-Body Physics. *PRX Quantum* **2**, 040201. <https://link.aps.org/doi/10.1103/PRXQuantum.2.040201> (4 Nov. 2021).
5. Casert, C., Vieijra, T., Whitelam, S. & Tamblyn, I. Dynamical Large Deviations of Two-Dimensional Kinetically Constrained Models Using a Neural-Network State Ansatz. *Phys. Rev. Lett.* **127**, 120602. <https://link.aps.org/doi/10.1103/PhysRevLett.127.120602> (12 Sept. 2021).
6. Hibat-Allah, M., Inack, E. M., Wiersema, R., Melko, R. G. & Carrasquilla, J. Variational neural annealing. *Nature Machine Intelligence* **3**, 952–961. ISSN: 2522-5839. <https://doi.org/10.1038/s42256-021-00401-3> (Nov. 2021).
7. Luo, D. *et al.* *Gauge Invariant Autoregressive Neural Networks for Quantum Lattice Models* 2021. <https://arxiv.org/abs/2101.07243>.
8. Morawetz, S., Vlugt, I. J. S. D., Carrasquilla, J. & Melko, R. G. U(1)-symmetric recurrent neural networks for quantum state reconstruction. *Physical Review A* **104**. <https://doi.org/10.1103/PhysRevA.104.012401> (July 2021).
9. Czischek, S., Moss, M. S., Radzihovsky, M., Merali, E. & Melko, R. G. *Data-Enhanced Variational Monte Carlo for Rydberg Atom Arrays* 2022. <https://arxiv.org/abs/2203.04988>.

10. Hibat-Allah, M., Melko, R. G. & Carrasquilla, J. Supplementing Recurrent Neural Network Wave Functions with Symmetry and Annealing to Improve Accuracy. *Machine Learning and the Physical Sciences (NeurIPS 2021)*. https://ml4physicalsciences.github.io/2021/files/NeurIPS_ML4PS_2021_92.pdf.